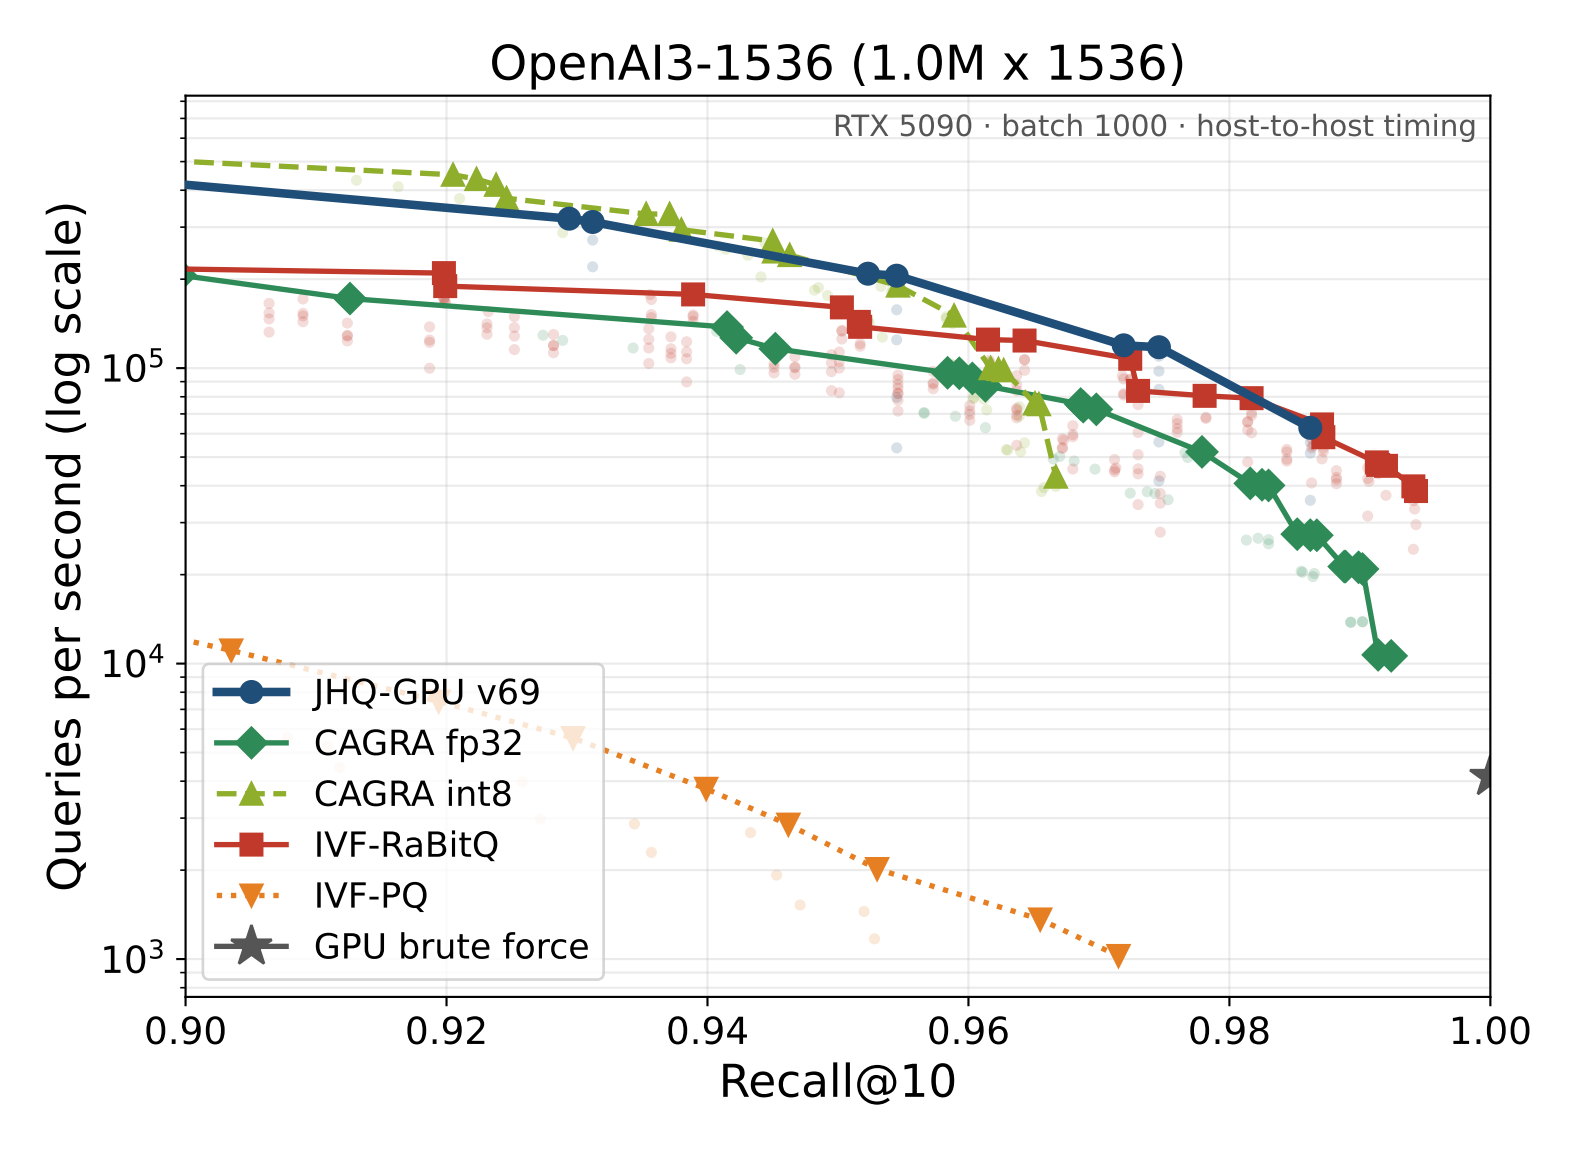
OpenAI3-1536 (1.0M x 1536)
RTX 5090 · batch 1000 · host-to-host timing
Queries per second (log scale)
10^5
10^4
10^3
JHQ-GPU v69
CAGRA fp32
CAGRA int8
IVF-RaBitQ
IVF-PQ
GPU brute force
0.90
0.92
0.94
0.96
0.98
1.00
Recall@10